

CS310 Natural Language Processing - Assignment 1: Neural Text Classification

Total points: 50

Train a neural network-based text classification model on the Chinese humor detection dataset. The dataset is collected by this project: <https://github.com/LaVi-Lab/CLEVA>.

Submit: The modified notebook file A1.ipynb and any dependent Python files.

1. Data Processing (15 points)

The original dataset (`train.jsonl` and `test.jsonl`) properly so that they can be loaded with `torch.utils.data.DataLoader`.

The raw `jsonl` data looks like this:

```
{"sentence": "和平说：你不挨屋儿做功课你跟这儿晃悠什么呢你", "choices": ["0", "1"],  
"label": [0], "id": "dev_292"} This is a negative (0) example (not humor)  
{"sentence": "季春生说：你还别说，就剩这部分机能还正常", "choices": ["0", "1"], "label":  
[1], "id": "dev_78"} This is a positive (1) example (humor)
```

Requirement:

- First, use a basic tokenizer that treats each single Chinese character (字) as a token, and discard all the non-Chinese tokens (English letters, numbers, punctuations).
- Incrementally improve your tokenizer by recognizing special patterns, such as consecutive digits, English words, and punctuations.

2. Build the Model (15 points)

Build the neural network model using `torch.nn` module.

Requirement:

- Use the bag-of-words method provided by `nn.EmbeddingBag`.
- Use at least 2 hidden layers when defining the fully-connected component. You can use the `torch.nn.Sequential` module for convenient implementation.

3. Train and Evaluate (10 points)

Make sure your train and evaluate code can run correctly and print the accuracy information properly.

Requirement:

- Train with sufficient epoch numbers (for example, 10) until you observe a good performance.
- Report the final test accuracy, precision, recall, and F-1 score.

4. Explore Word Segmentation (10 points)

Include word segmentation program to the processing step and observe how that will affect the classification performance. You can think of word segmentation as a more advanced tokenizer, which can group multiple characters (字) to a word (词). Note that the vocabulary needs be re-built once segmentation is applied, whose size will get smaller.

Requirement:

- Install a word segmentation package, such as `jieba`: <https://github.com/fxsjy/jieba>. Use it to process the data and save as the new segmented data.
- Run through the train/evaluate process on the segmented data and compare the performance (accuracy and F1) with the original data.