

# CS310 Natural Language Processing

## 自然语言处理

### Lecture 12 - Question Answering and Group Project Discussion

Instructor: Yang Xu

主讲人：徐炆

xuyang@sustech.edu.cn

# Overview

- **Question Answering (QA)**
  - What is QA?
  - Information Retrieval; Tf-idf
  - Retriever-based QA; Datasets
    - Answer Span Extraction
    - Retrieval-Augmented Generation
- Project Discussion

# What is Question Answering?

- To build a system that **automatically** answer questions posed by human in natural language

“The Ultimate Question  
Of Life, The Universe,  
and Everything”



(from movie *Hitchhiker's Guide to the Galaxy*)

# QA System focuses on *factoid* questions

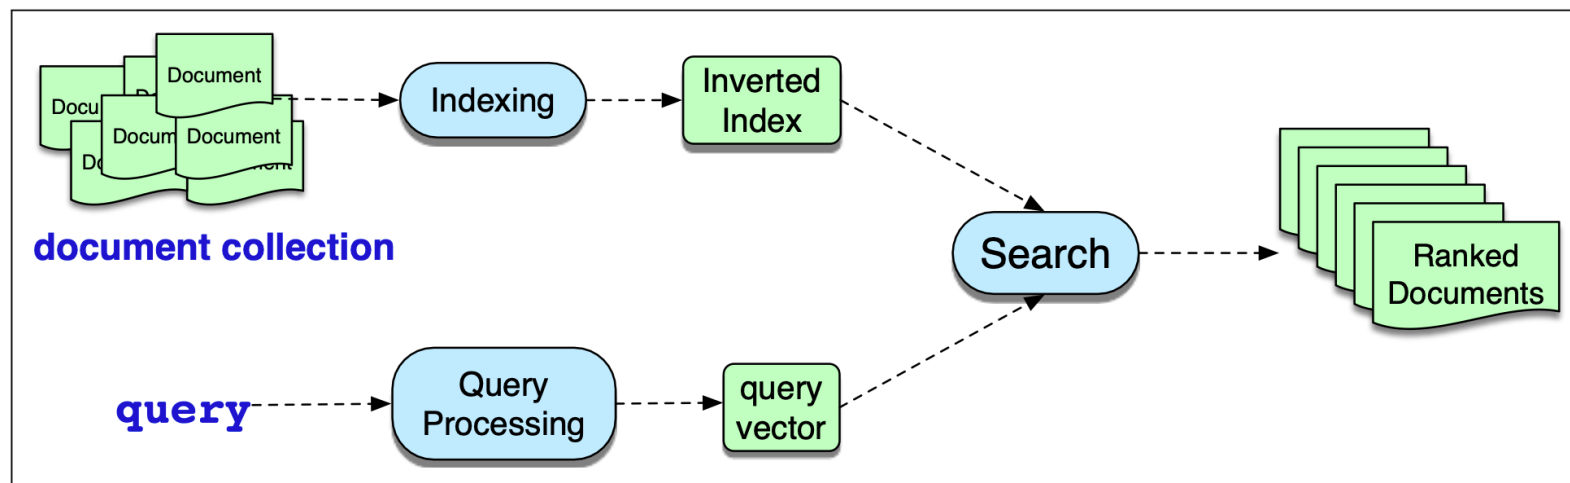
- **factoid questions:** Questions that can be answered with simple facts expressed in short texts
- Ex. 1: [Where is the Louvre Museum located?](#)
- Ex. 2: [What is the distance from Moon to Earth?](#)
- One way: to directly ask a large language model (LLM)
  - Using prompts: “Q: [What is the distance from Moon to Earth?](#) A: ”
- **Problems:**
- LLMs hallucinate; not calibrated
- No access to proprietary/private/personal data: email, private documents, ...

# Current Solution to QA

- Two-stage **retriever/reader** model
  - Stage 1. Retriever algorithms: Use information retrieval (IR) to retrieve relevant documents
  - Stage 2. Reader algorithms: Either **extract** or **generate** an answer

# Brief Overview of Information Retrieval (IR)

- **Information retrieval, IR:** Retrieval of all kinds of media based on user information needs. IR system  $\approx$  **search engine**
- We focus on **ad hoc (临时) retrieval**: a user poses a query to an IR system, which then returns an ordered set of documents from some collection



**Figure 14.1** The architecture of an ad hoc IR system.

**Query:** a user's information need expressed as a set of **terms**

**Term** refers to a word/phrase in a collection of documents

Figure from SLP3, Ch 14

# How to match a document a query?

- Compute a term weight for each document term
- Common method: **tf-idf** and BM25
  - **tf**: term frequency
  - **idf**: inverse document frequency
- $\text{tf-idf} \triangleq \text{tf} \times \text{idf}$  (product of two)

term  $t$ ; document  $d$

$$\text{tf}_{t,d} = \begin{cases} 1 + \log_{10} \text{count}(t, d) & \text{if } \text{count}(t, d) > 0 \\ 0 & \text{otherwise} \end{cases}$$

- **tf**: words that occur more often in a document are likely to be informative about the document's content
- Use the  $\log_{10}$  of word frequency count rather than raw count
- Why? A word appearing 100 times doesn't make it 100 times more likely

# Tf-idf

$$\text{tf}_{t,d} = \begin{cases} 1 + \log_{10} \text{count}(t,d) & \text{if } \text{count}(t,d) > 0 \\ 0 & \text{otherwise} \end{cases}$$

term  $t$ ; document  $d$

term occurs 0 times in document:  $\text{tf} = 0$   
term occurs 1 times in document:  $\text{tf} = 1$   
term occurs 10 times in document:  $\text{tf} = 2, \dots$

- **document frequency**  $\text{df}_t$  of a term  $t$  is the number of documents it occurs in
- Terms that occur in only **a few** documents are useful for discriminating those documents from the rest of the collection;
- terms that occur across the entire collection aren't as helpful (*the, a, an, ...*)
- **inverse document frequency** or **idf** is defined as:

$$\text{idf}_t = \log_{10} \frac{N}{\text{df}_t}$$

$N$ : total number of documents  
The fewer documents in which  $t$  occurs, the higher  $\text{idf}_t$



# Inverse document frequency example

- Some idf values for some words in the corpus of Shakespeare plays

Word	df	idf
Romeo	1	1.57
salad	2	1.27
Falstaff	4	0.967
forest	12	0.489
battle	21	0.246
wit	34	0.037
fool	36	0.012
good	37	0
sweet	37	0

Extremely informative words that occur in only one play like *Romeo*

*good* or *sweet* are completely non-discriminative since they occur in all 37 plays

# Scoring with tf-idf

- We can score document  $d$  by the cosine of its vector  $\vec{d}$  with the query vector  $\vec{q}$ :

$$\text{score}(q, d) = \cos(\vec{q}, \vec{d}) = \frac{\vec{q} \cdot \vec{d}}{|\vec{q}| \cdot |\vec{d}|}$$

- in which  $\vec{q}$  and  $\vec{d}$  are vectors of query length  $n$ , whose values are the **tf-idf** values (normalized):

$$\begin{aligned} \vec{q} &= \frac{[\text{tf-idf}(t_1, q), \dots, \text{tf-idf}(t_n, q)]}{\sqrt{\sum_{t \in q} \text{tf-idf}^2(t, q)}} \\ \vec{d} &= \frac{[\text{tf-idf}(t_1, d), \dots, \text{tf-idf}(t_n, d)]}{\sqrt{\sum_{t \in d} \text{tf-idf}^2(t, d)}} \end{aligned} \quad \Rightarrow \quad \text{score}(q, d) = \sum_{t_i \in q} \frac{\text{tf-idf}(t_i, q)}{\sqrt{\sum_{t \in q} \text{tf-idf}^2(t, q)}} \cdot \frac{\text{tf-idf}(t_i, d)}{\sqrt{\sum_{t \in d} \text{tf-idf}^2(t, d)}}$$

# Tf-idf scoring example

- A collection of 4 nano documents

**Query:** sweet love

**Doc 1:** Sweet sweet nurse! Love?

**Doc 2:** Sweet sorrow

**Doc 3:** How sweet is love?

**Doc 4:** Nurse!

Using a variant of tf-idf cosine score, by dropping the idf term for the document (for better perf.)

$$\text{score}(q, d) = \sum_{t \in q} \frac{\text{tf}_{t,q} \cdot \text{idf}_t}{\sqrt{\sum_{q_i \in q} \text{tf-idf}^2(q_i, q)}} \cdot \frac{\text{tf}_{t,d} \cdot \text{idf}_t}{\sqrt{\sum_{d_i \in d} \text{tf-idf}^2(d_i, d)}}$$

Query vector  $\vec{q} = (0.383, 0.924)$

Query						
word	cnt	tf	df	idf	tf-idf	n'lized = tf-idf/ q
sweet	1	1	3	0.125	0.125	0.383
nurse	0	0	2	0.301	0	0
love	1	1	2	0.301	0.301	0.924
how	0	0	1	0.602	0	0
sorrow	0	0	1	0.602	0	0
is	0	0	1	0.602	0	0
$ q  = \sqrt{.125^2 + .301^2} = .326$						

# Tf-idf scoring example

Query vector  $\vec{q} = (0.383, 0.924)$

Document 1					
word	cnt	tf	tf-idf	n'lized	$\times q$
sweet	2	1.301	0.163	0.357	0.137
nurse	1	1.000	0.301	0.661	0
love	1	1.000	0.301	0.661	0.610
how	0	0	0	0	0
sorrow	0	0	0	0	0
is	0	0	0	0	0
$ d_1  = \sqrt{.163^2 + .301^2 + .301^2} = .456$					

$$\vec{d}_1 = (0.357, 0.661)$$

$$\text{score}(\vec{q}, \vec{d}_1) = 0.747$$

Therefore,  $d_1$  is more relevant

Document 2					
word	cnt	tf	tf-idf	n'lized	$\times q$
sweet	1	1.000	0.125	0.203	0.0779
nurse	0	0	0	0	0
love	0	0	0	0	0
how	0	0	0	0	0
sorrow	1	1.000	0.602	0.979	0
is	0	0	0	0	0
$ d_2  = \sqrt{.125^2 + .602^2} = .615$					

$$\vec{d}_2 = (0.203)$$

$$\text{score}(\vec{q}, \vec{d}_1) = 0.0779$$

**Query:** sweet love

**Doc 1:** Sweet sweet nurse! Love?

**Doc 2:** Sweet sorrow

# Efficient Implementation: Inverted Index

- The basic search problem in IR is to find all documents  $d \in \mathcal{C}$  that contain a term  $q \in \mathcal{Q}$
- Use the data structure **inverted index**: given a query term, returns a list of documents that contain the term
- Contains two parts: *dictionary* and *postings*

**dictionary**: a list of terms, each pointing to a postings list for the term (including document frequency)

how {1}	→	3 [1]
is {1}	→	3 [1]
love {2}	→	1 [1] → 3 [1]
nurse {2}	→	1 [1] → 4 [1]
sorry {1}	→	2 [1]
sweet {3}	→	1 [2] → 2 [1] → 3 [1]

**posting list**: a list of document IDs associated with each term (including term frequency etc.)

# IR with Dense Vectors

- **Flaws of TF-IDF -- Vocabulary mismatch problem:** it only works if there is exact overlap of words between the query and document
- **Solution:** Using dense vectors to represent queries/documents  
dating back to Latent semantic indexing vectors, all the way to modern times via encoders like BERT

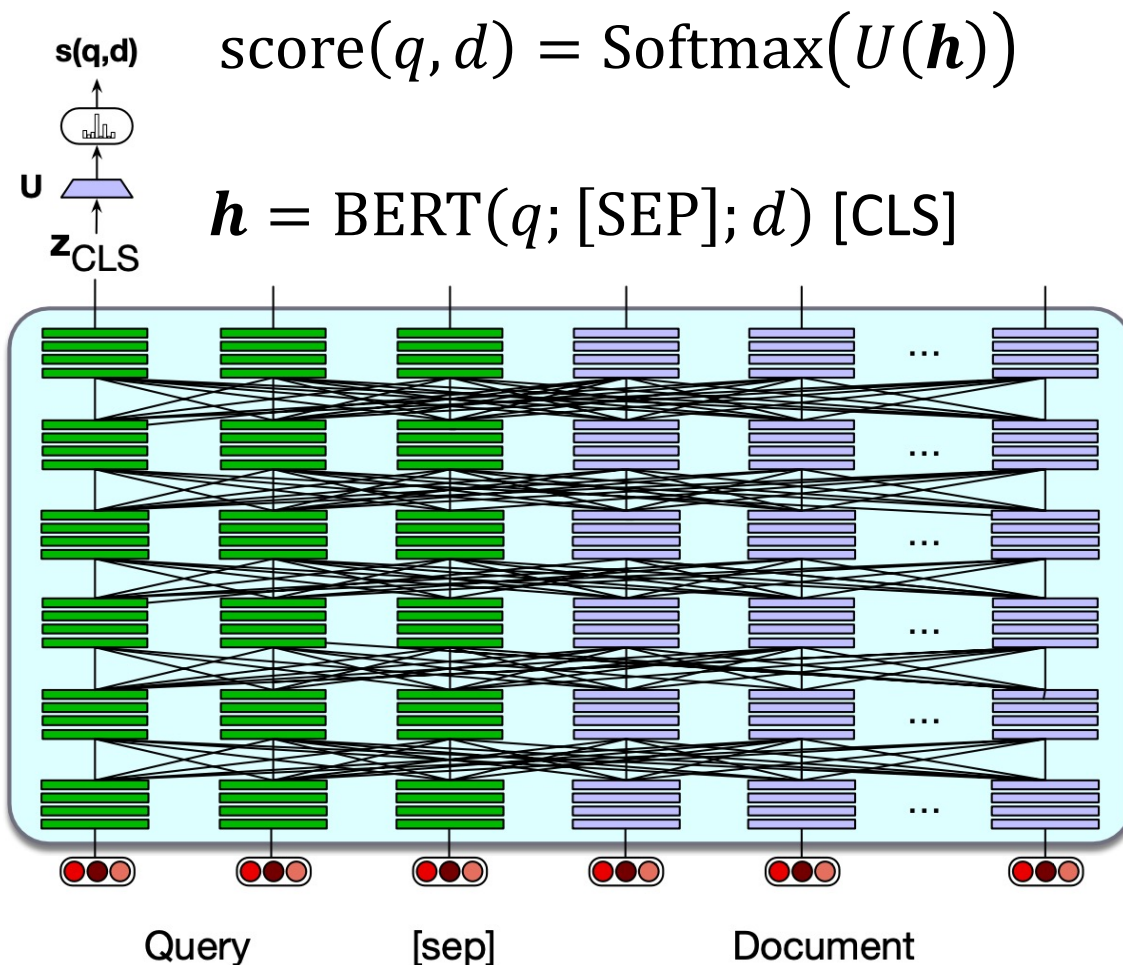
Present both query  $q$  and document  $d$  to a single encoder, allowing self-attention to see all tokens from both  $q$  and  $d$

$$\mathbf{h} = \text{BERT}(q; [\text{SEP}]; d) [\text{CLS}]$$

$$\text{score}(q, d) = \text{Softmax}(U(\mathbf{h}))$$

Predict the similarity score between  $q$  and  $d$

# Single BERT Encoder for IR



In practice, documents are broken up into smaller passages such as non-overlapping fixed-length chunks of  $\sim 100$  tokens,

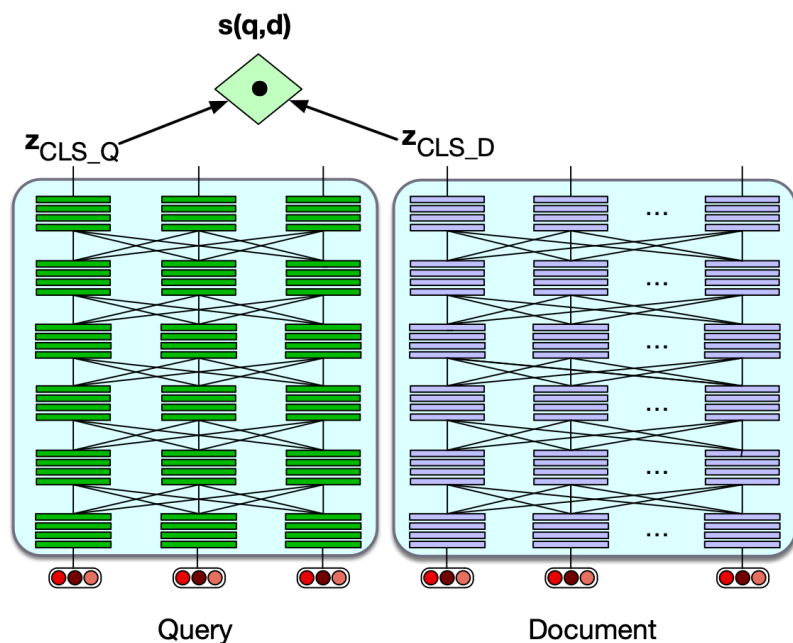
so that the  $q$  and  $d$  can fit in the BERT 512-token window

**Drawback: expense in computation!**

Every time we get a query, have to pass every single document through a BERT encoder jointly with the new query!

# More Efficient Way: Bi-Encoder

- Two separate encoder models:  
one to encode the query  $\text{BERT}_Q$ , and one to encode the document,  $\text{BERT}_D$
- Encode each document and store the document vectors in advance
- When a query comes in, just encode this query, and compute the dot product between it and each candidate document



$$\mathbf{h}_q = \text{BERT}_Q(q) [\text{CLS}]$$

$$\mathbf{h}_d = \text{BERT}_D(d) [\text{CLS}]$$

$$\text{score}(q, d) = \mathbf{h}_q \cdot \mathbf{h}_d$$

Cheaper in computation, but less accurate, since it does not take full advantage of the interaction between query tokens and document tokens

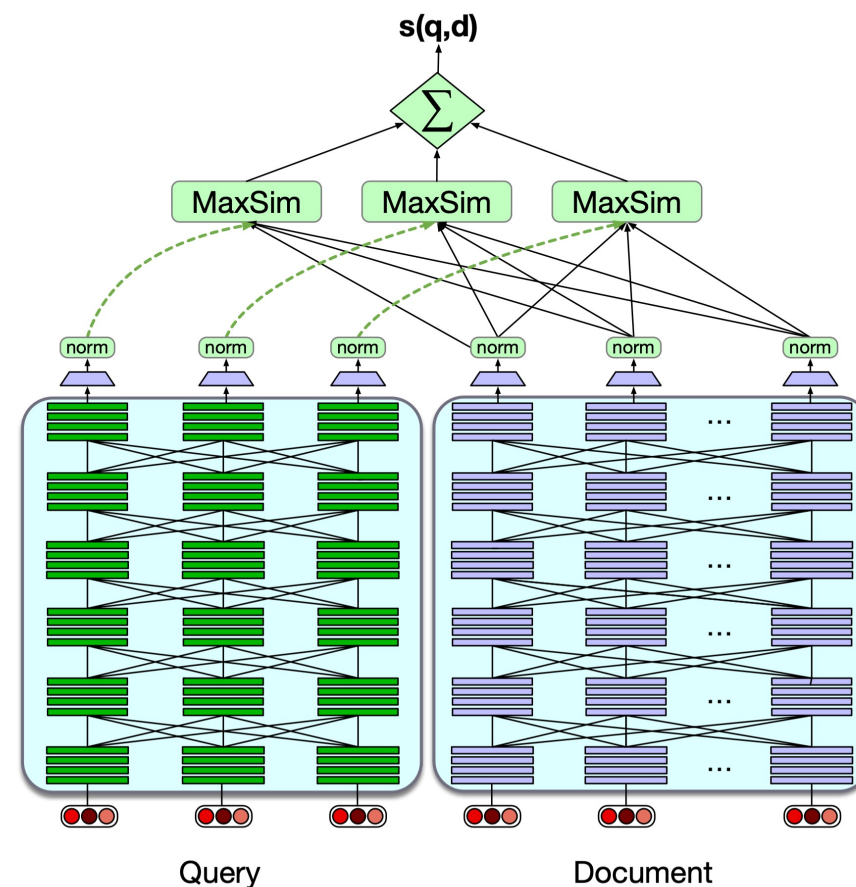


# Alternative: Token-level similarity scores

- **ColBERT** (Khattab et al., 2021) computes the score between a query  $q$  and a document  $d$  as a sum of maximum similarity (MaxSim) between tokens in  $q$  and tokens in  $d$

$$\text{score}(q, d) = \sum_{i=1}^N \max_{j=1}^m \mathbf{E}_{q_i} \cdot \mathbf{E}_{d_j}$$

More accurate than the bi-encoder method



# Implementation Efficiency

- For dense vector-based IR, efficiency is also an important issue
- since every possible document must be ranked for its similarity to the query
- Bottle-neck: Finding the set of document vectors that have the highest dot product with a query vector -- **nearest neighbor search** problem
- Can be approximated with algorithms like Faiss (Johnson et al., 2017)

# Current Solution to QA

- Two-stage **retriever/reader** model
- Stage 1. Retriever algorithms: Use information retrieval (IR) to retrieve relevant documents
- Stage 2. Reader algorithms: Either extract or generate an answer

Reader

- Extractor: **span extraction**  $\Rightarrow$  find spans of text that answer the question over the retrieved passages
- Generator: **retrieval-augmented generation**  $\Rightarrow$  Take a large pretrained LM, design the prompt based on the retrieved passage, and generate the answer token by token

# Two-staged retriever/reader model

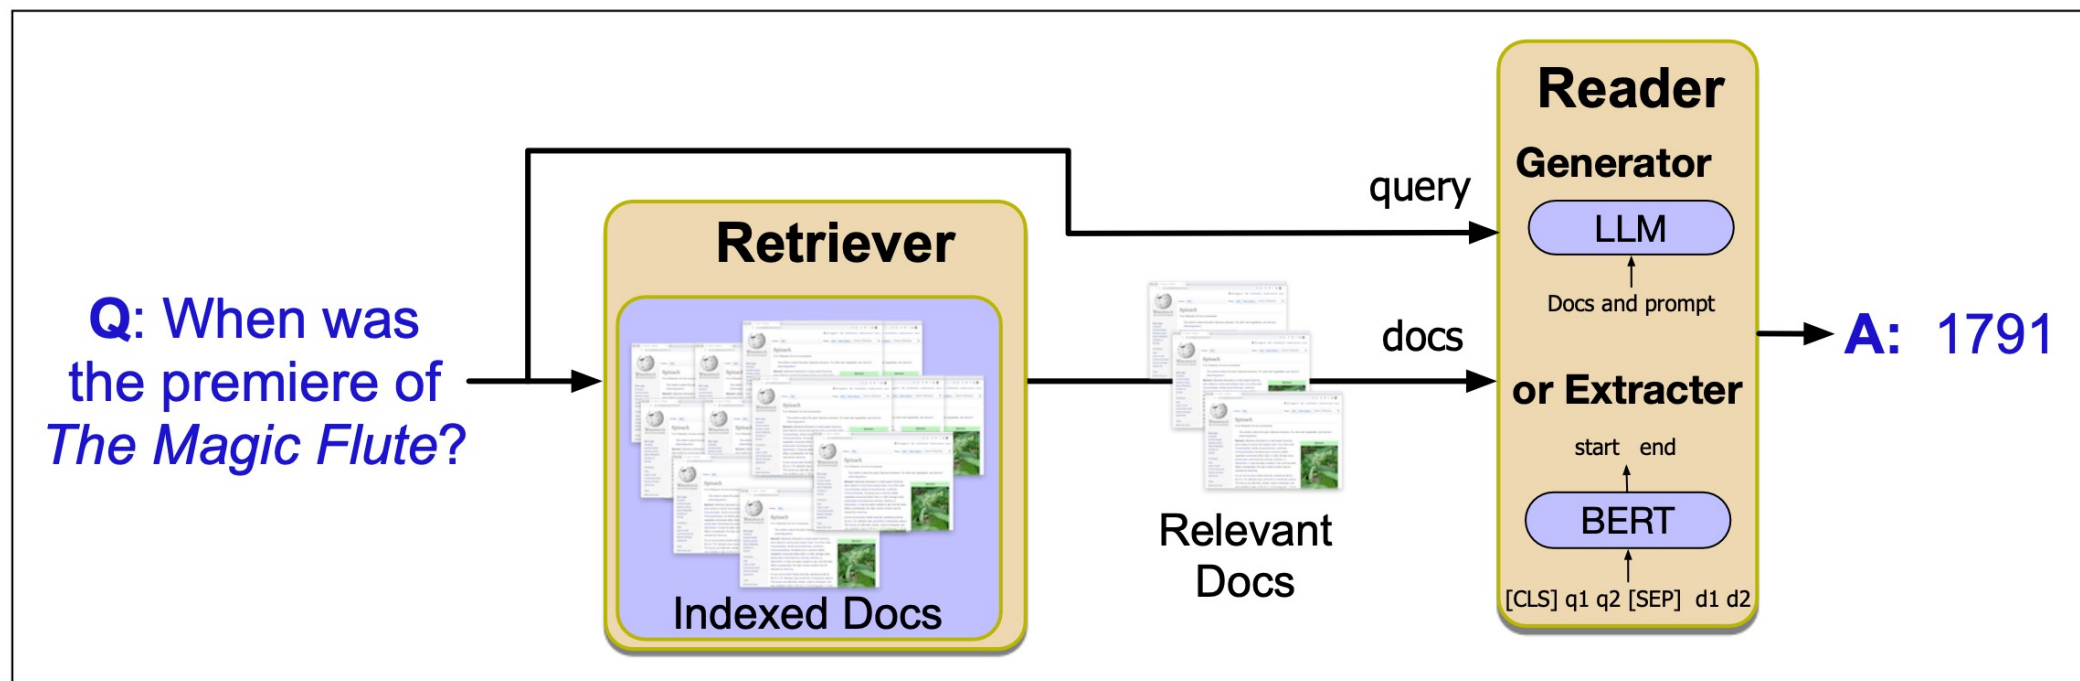


Figure from SLP3, Ch 14

# Reader: Answer Span Extraction

Beyoncé Giselle Knowles-Carter (born September 4, 1981) is an American singer, songwriter, record producer and actress. Born and raised in **Houston, Texas**, she performed in various **singing and dancing** competitions as a child, and rose to fame in the late 1990s as lead singer of R&B girl-group Destiny's Child. Managed by her father, Mathew Knowles, the group became one of the world's best-selling girl groups of all time. Their hiatus saw the release of Beyoncé's debut album, *Dangerously in Love* (**2003**), which established her as a solo artist worldwide, earned five Grammy Awards and featured the Billboard Hot 100 number-one singles "Crazy in Love" and "Baby Boy".

Q: "In what city and state did Beyoncé grow up?"

A: "**Houston, Texas**"

Q: "What areas did Beyoncé compete in when she was growing up?"

A: "**singing and dancing**"

Q: "When did Beyoncé release *Dangerously in Love*?"

A: "**2003**"

**Figure 14.11** A (Wikipedia) passage from the SQuAD 2.0 dataset ([Rajpurkar et al., 2018](#)) with 3 sample questions and the labeled answer spans.

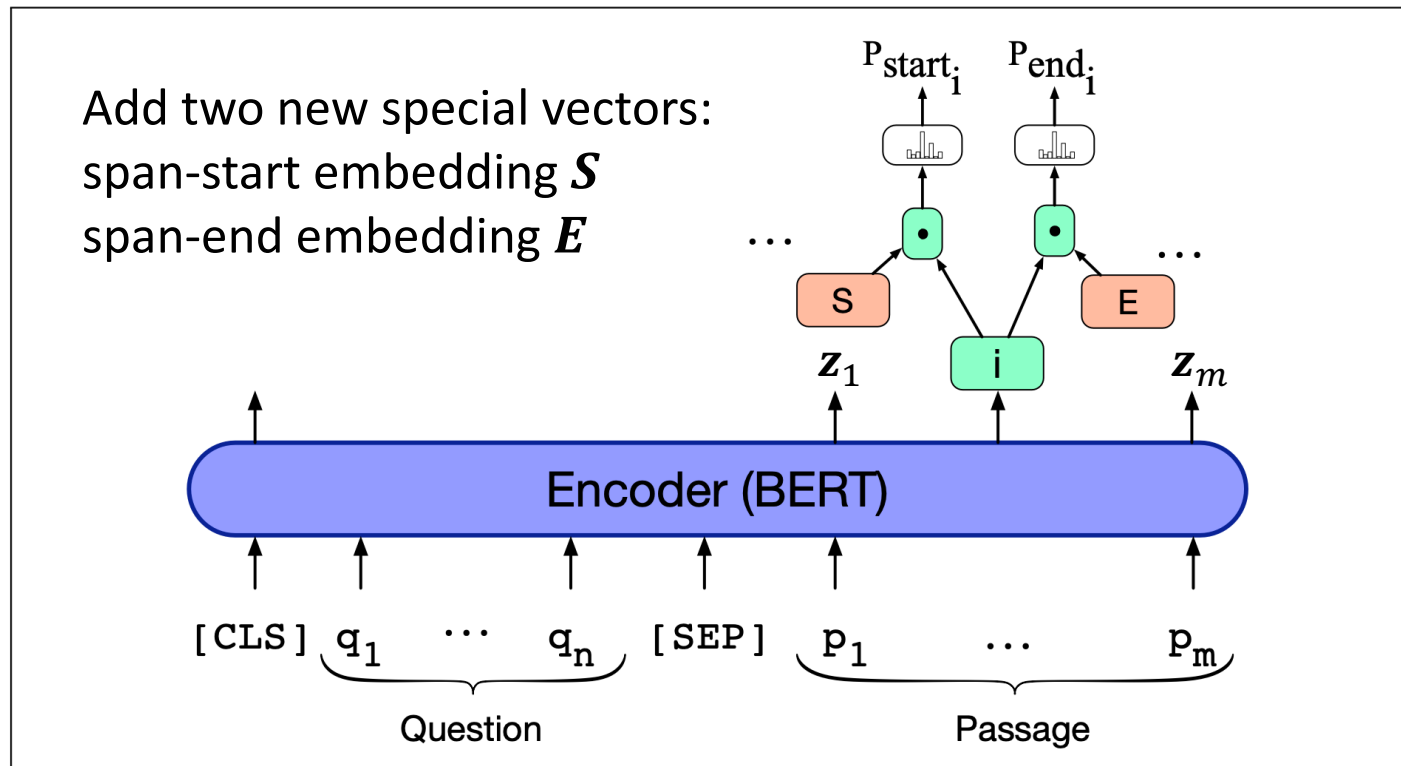
**Span labeling** task: identify in the passage a span (continuous string of text) that constitutes an answer

# Span Labeling

- Given a question  $q$  of  $n$  tokens  $q_1, \dots, q_n$  and a passage  $p$  of  $m$  tokens  $p_1, \dots, p_m$
- **Goal:** compute the probability  $P(a|q, p)$  of each possible span  $a$  is the answer
- Span  $a$  starts at position  $a_s$  and ends at  $a_e$ , then estimate the probability by:
- $P(a|q, p) = P_{\text{start}}(a_s|q, p)P_{\text{end}}(a_e|q, p)$
- For each token  $p_i$  in passage, compute two probabilities:
  - $P_{\text{start}}(i) \Rightarrow p_i$  is the start of answer span
  - $P_{\text{end}}(i) \Rightarrow p_i$  is the end of answer span
- Goal becomes:

$$\max_{i, j \in [1, m]; j \geq i} P_{\text{start}}(i)P_{\text{end}}(j)$$

# Span Labeling



$$P_{start}(i) = \frac{\exp(S \cdot z_i)}{\sum_j \exp(S \cdot z_j)}$$

$$P_{end}(i) = \frac{\exp(E \cdot z_i)}{\sum_j \exp(E \cdot z_j)}$$

Goal:  $\max_{i,j \in [1,m]; j \geq i} P_{start}(i) P_{end}(j)$



$$\max_{i,j \in [1,m]; j \geq i} S \cdot z_i + E \cdot z_j$$

The score for candidate span  
from position  $i$  to  $j$

# What if answer is not contained in passage?

- Many datasets contain (question, passage) pairs in which the answer is not contained in the passage
- Need a way to estimate this “none” probability
- Done by treating [CLS] token as the answer, i.e.,  $a_s$  and  $a_e$  all point to [CLS]



# Retrieval-based QA Datasets

- Reading comprehension datasets containing tuples of (*passage*, *question*, *answer*)
- Including *passage* eliminates the need for information retrieval
- A system can be trained to predict a span in passage as answer, given a question
- Stanford Question Answering Dataset (**SQuAD**) (Rajpurkar et al., 2016)
  - Over 150,000 questions
  - Passage from Wikipedia; SQuAD 2.0 includes unanswerable questions

# Retrieval-based QA Datasets

- **HotpotQA** dataset (Yang et al., 2018): Showing crowd workers multiple context documents and asked to create questions that require reasoning
- Both SQuAD and HotpotQA are created by annotators who have first read the passage may make their questions easier to answer
- Datasets from questions that were not written with a passage in mind
- **TriviaQA** dataset (Joshi et al., 2017) Trivia: 琐事, 娱乐和消遣
  - 94K questions written by trivia enthusiasts, with supporting documents (Wikipedia and web)
  - 650K question-answer-evidence triples
  - Relatively complex, compositional questions
  - Requires more cross sentence reasoning

# Retrieval-based QA Datasets

- **MS MARCO** (Microsoft Machine Reading Comprehension) (Nguyen et al., 2016)
  - 1 million real anonymized questions from Microsoft Bing query logs
  - with a human generated answer and 9 million passages
  - <https://microsoft.github.io/msmarco/>
- **Natural Questions** dataset (Kwiatkowski et al., 2019)
  - Anonymized queries to the Google search engine
  - Annotators are presented a query, along with a Wikipedia page from the top 5 search results
  - To annotate a paragraph-length **long** answer and a **short** span answer, or mark **null** if the text doesn't contain the paragraph.
  - <https://ai.google.com/research/NaturalQuestions>

# Current Solution to QA

- Two-stage **retriever/reader** model
- Stage 1. Retriever algorithms: Use information retrieval (IR) to retrieve relevant documents
- Stage 2. Reader algorithms: Either extract or generate an answer

Reader

- Extractor: **span extraction**  $\Rightarrow$  find spans of text that answer the question over the retrieved passages
- Generator: **retrieval-augmented generation**  $\Rightarrow$  Take a large pretrained LM, design the prompt based on the retrieved passage, and generate the answer token by token

# Reader: Retrieval-Augmented Generation (RAG)

- Cast the QA task as word prediction: feeding the LM a question and a token like “A: ” -- suggesting the answer should come next

Q: Who wrote the book ‘‘The Origin of Species’’? A:  $x_1, x_2, \dots, x_n$

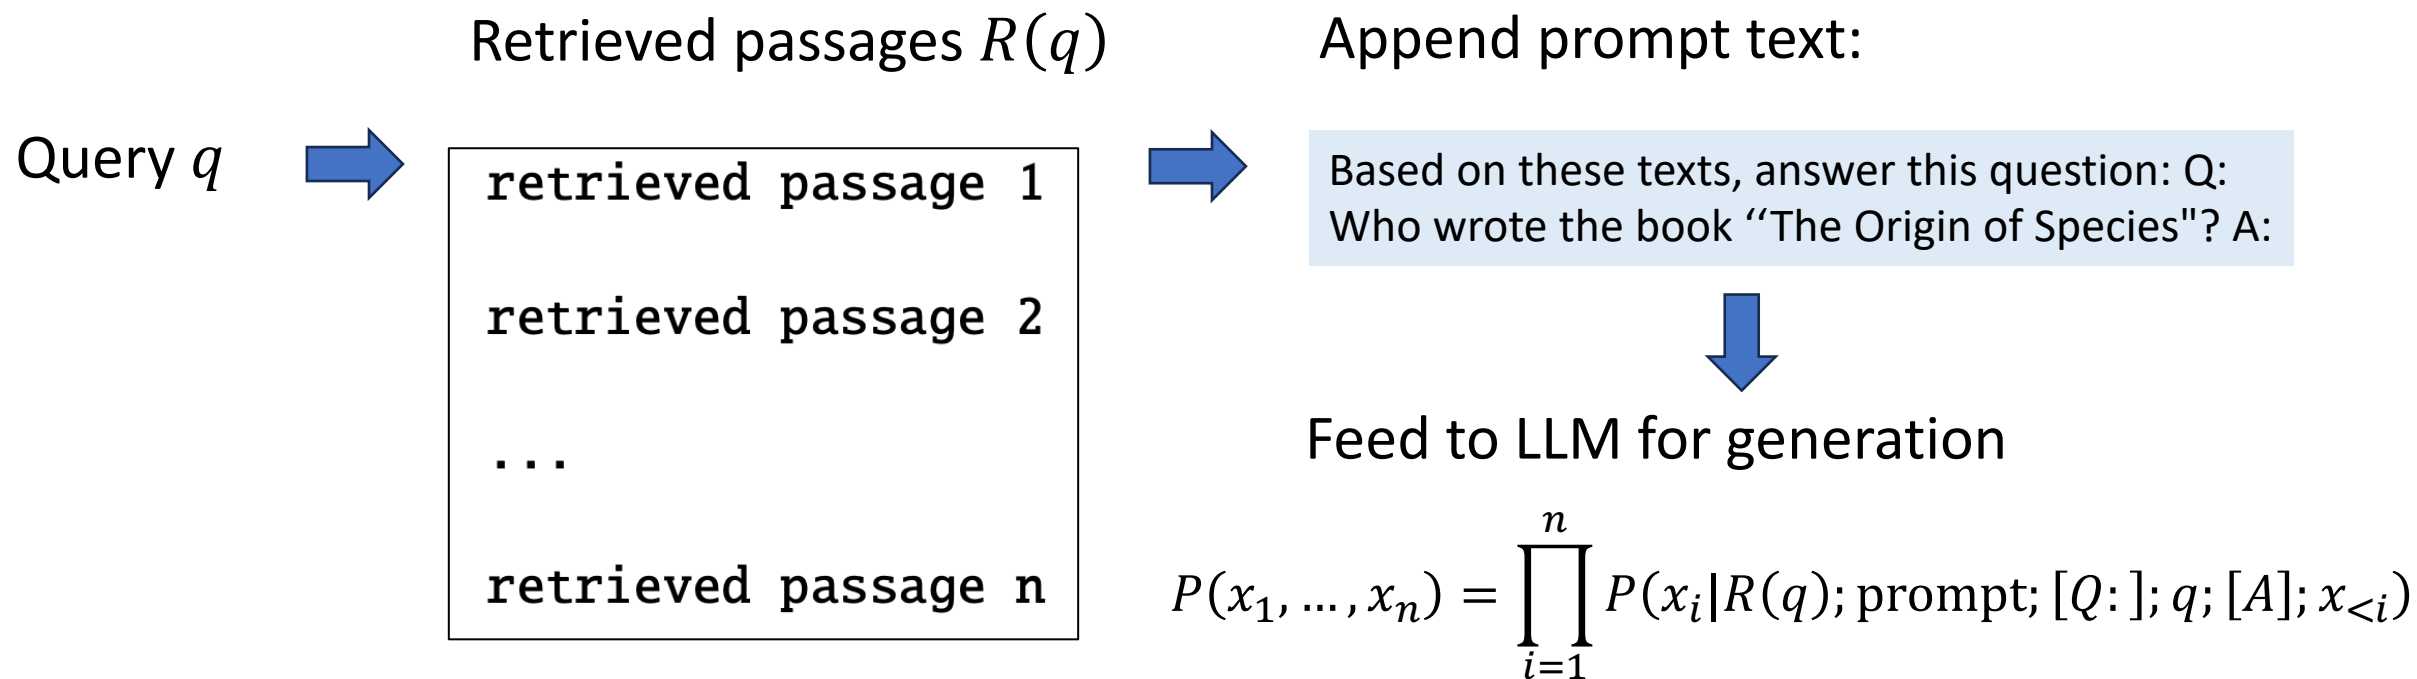
[Q:]  $\underbrace{\hspace{15em}}_q$  [A:]

Conditional generation that optimizes:  $P(x_1, \dots, x_n) = \prod_{i=1}^n P(x_i | [Q: ]; q; [A]; x_{<i})$

**Problems:** hallucination; no access to proprietary data

# Retrieval-Augmented Generation (RAG)

- Idea: Conditioning on the retrieved passages as part of the prefix perhaps with some prompt like “Based on these text, answer this question”



# RAG details

- Just like span-based extractor, RAG requires a successful retriever, in two-stage setting as well
- **Multi-hop** architecture may be needed: a query  $q$  is used to retrieve documents, which are then appended to original  $q$  for a *second* stage retrieval
- Detailed prompt engineering is needed
- When combining private data with public data, externally hosted LLMs may be concerned

# Evaluation of Retrieval-based QA

- QA is commonly evaluated using **mean reciprocal rank**, or **MRR**

$$\text{MRR} = \frac{1}{|Q|} \sum_{i=1}^{|Q|} \frac{1}{\text{rank}_i}$$

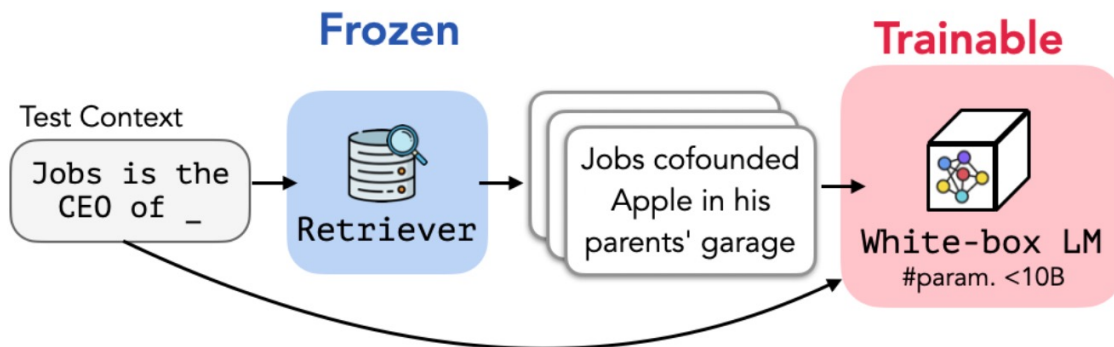
- If the system returned 5 answers, but the first 3 are wrong, then the highest-ranked correct answer is ranked 4<sup>th</sup> and thus the reciprocal rank is  $\frac{1}{4}$
- Alternative methods:
- **Exact match**: The % of predicted answers that match the gold answer exactly
- **F-1 score**: average F-1 over all questions



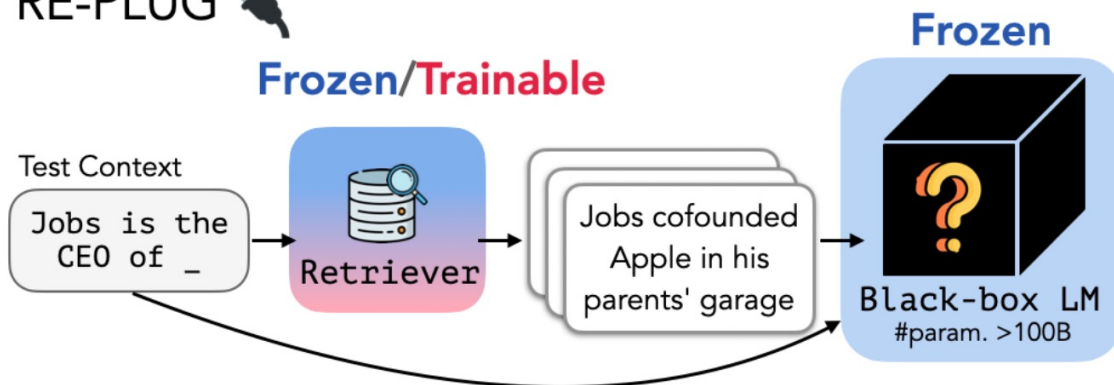
# RAG Example: RePLUG

Shi et al., 2023

Previous



RE-PLUG

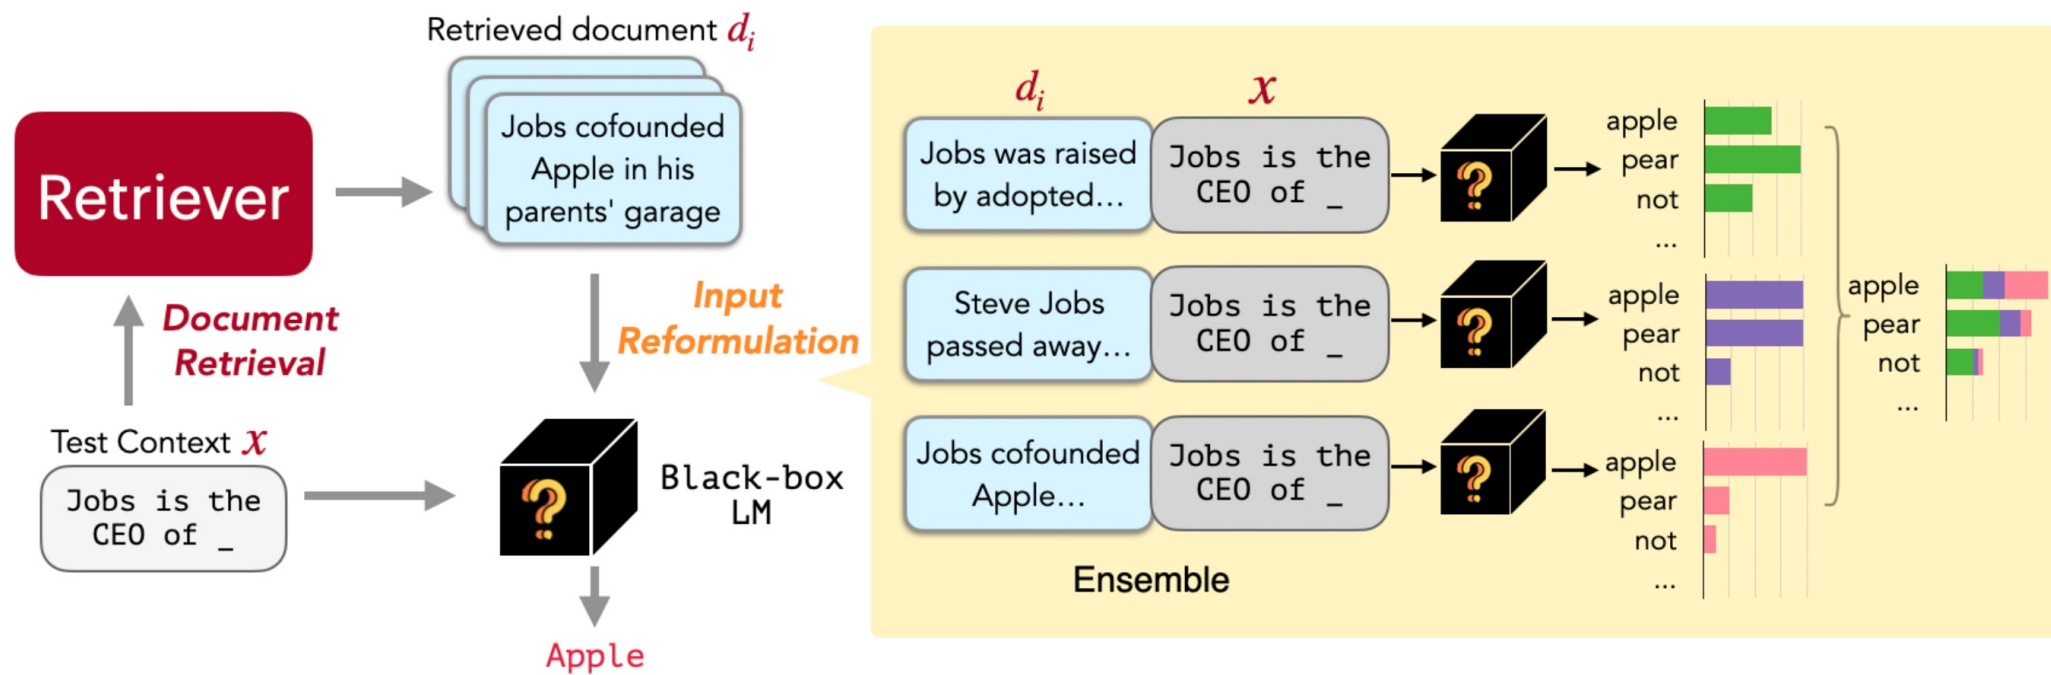


Treats the LM as a black box and augments it with a tuneable retrieval model

- Simply prepends retrieved documents to the input
- LM can be used to supervise the retrieval model

# RAG Example: RePLUG

Shi et al., 2023



- Given an input context, REPLUG first retrieves a small set of relevant documents from an external corpus
- Then it prepends each document separately to the input context and ensembles output probabilities from different passes

# Overview

- Question Answering (QA)
- **Project Discussion**
  - Do and Do-Nots
  - Trending Topics
  - Benchmarks
    - TriviaQA, Natural Questions, HotpotQA
    - GLUE, SuperGLUE

# Project: What to do?

- Default project: **BERT + Fine-tuning on downstream tasks**

- Examples:

- BERT + Sequence Classification

- Sentiment classification
- Paraphrase detection
- Semantic similarity etc.

Code template  
provided

- Or, BERT + QA on SQuAD, TriviaQA, Natural Questions etc.
- Or, BERT + Translation



# BERT + Sequence Classification

Source code credit to: Stanford  
CS 224N Winter 2023 Default  
Final Project: Multitask BERT  
<https://github.com/gpoesia/minbert-default-final-project>

- Primary Task: **Sentiment classification**
- Training dataset: Stanford Sentiment Treebank (SST) on movies
  - Train: 8545 lines of (sentence, score) pairs; score from 1 (neg) to 4 (pos)
  - Dev: 1102 lines
- Requirement
  - Finish the implementation of BERT ([bert.py](#), **skeleton provided, with six TODOs**); Initialized from pretrained model
  - Fine-tune it on SST data ([classifier.py](#), **mostly implemented with two TODOs**)
  - Extend and improve it in various ways:
    - Multi-task task through **paraphrase detection** and **semantic similarity regression** tasks ([multitask\\_classifier.py](#), **three new TODOs**)
    - Different tasks correspond to different [predict\\_xxx\(\)](#) functions in forward function

# What to do with custom projects

- If you:
  - Have some research project that you're excited about (and are possibly already working on)
  - You want to try to do something different
  - You want to see more of the process of defining a research goal, finding data and tools, and working out something you could do that is interesting, and how to evaluate it
- Then: Do the custom final project
- **Requirement:** must substantively involves both human language and neural networks

# Project: What not to do?

- Train BIG models from scratch
  - Be realistic about the scale of compute you can do
  - You do not have the resources to train your own GPT-2 model from scratch
  - You probably do not have the resources to load a 7- to 11-B model (Llama-2, ChatGLM-3, Mistral-7B, T5-11B etc.)

# Some trending topics

- Evaluating and improving models for something other than accuracy
  - Adaptation when there is domain shift
  - Evaluating the robustness of models in general
- Empirical work looking at what large pre-trained models have learned
- Get knowledge and good task performance without much data
- Bias, trustworthiness, and interpretability of large models
- Low resource languages or problems

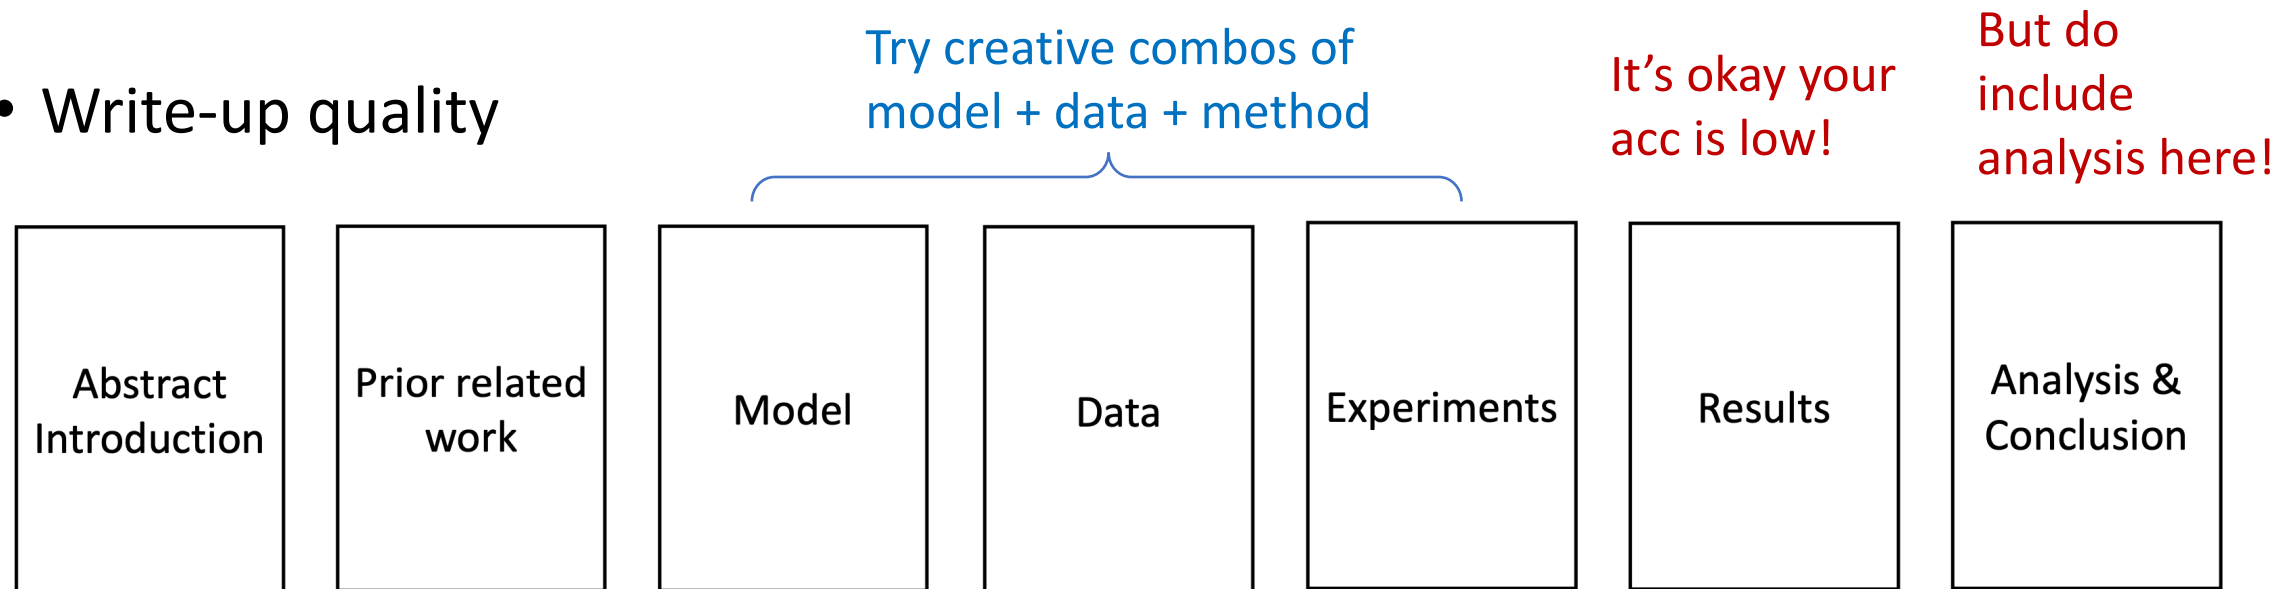


# Some trending topics

- Building small, performant models can be cool!
- Model pruning/quantization
  - QLoRA; Pruning; Compression:  
<https://proceedings.mlr.press/v119/li20m/li20m.pdf>;  
<https://arxiv.org/pdf/2004.07320>
  - Efficient Open-domain QA: <https://efficientqa.github.io/> (within 6GB mem)
- Baby LM challenge: <https://babylm.github.io/index.html>
  - Efforts on optimizing pretraining given data limitations inspired by human development
  - 100M to 10M word text data

# Grading: Project and Presentation

- Write-up quality



- Focus on what you have done  
-- not on the amazing ChatGPT output showing that “look, it works zero-shot”
- Minimal 5 pages (template provided)

# Important Dates

- In-class presentation of project: Week 16, Tuesday, June 4<sup>th</sup>, 2024
  - 7 minutes presentation + 3 minutes QA
- Project report due: Friday 11:59 PM, June 7<sup>th</sup>, 2024

# References

- Nguyen, T., Rosenberg, M., Song, X., Gao, J., Tiwary, S., Majumder, R., & Deng, L. (2016). Ms marco: A human-generated machine reading comprehension dataset.
- Chen, D., Fisch, A., Weston, J., & Bordes, A. (2017). Reading wikipedia to answer open-domain questions. *arXiv preprint arXiv:1704.00051*.
- Joshi, M., Choi, E., Weld, D. S., & Zettlemoyer, L. (2017). Triviaqa: A large scale distantly supervised challenge dataset for reading comprehension. *arXiv preprint arXiv:1705.03551*.
- Johnson, J., Douze, M., & Jégou, H. (2019). Billion-scale similarity search with GPUs. *IEEE Transactions on Big Data*, 7(3), 535-547.
- Karpukhin, V., Oğuz, B., Min, S., Lewis, P., Wu, L., Edunov, S., ... & Yih, W. T. (2020). Dense passage retrieval for open-domain question answering. *arXiv preprint arXiv:2004.04906*.
- Shi, W., Min, S., Yasunaga, M., Seo, M., James, R., Lewis, M., ... & Yih, W. T. (2023). Replug: Retrieval-augmented black-box language models. *arXiv preprint arXiv:2301.12652*.